

IN4010 Practical Assignment: Reinforcement Learning

November 16, 2021

Contents

1	Intro	2
2	Set up	2
3	Environments	3
3.1	Q-Learning	3
3.2	Deep Q-Learning	3
4	Q-Learning	4
5	Deep Q-Learning	5

1 Intro

In this practical assignment, you will develop some hands-on experience with implementing and applying reinforcement learning techniques. The assignment consists of two parts: one focusing on regular Q-learning, and one focusing on Q-learning with a *deep Q network* [1].

Coding Framework

In the zip file you can find the following files:

1. For the Q-learning part:
 - (a) `simple_grid.py`, the environment(s) we will be using, see Section 3 for a description.
 - (b) `q_learning_main.py`, contains a main loop you can use to run your experiments.
 - (c) `q_learning_skeleton.py`, file for the Q-learning agent, to be coded up.
2. For the deep Q-learning part:
 - (a) `deep_q_learning_main.py`, contains a main loop you can use to run your experiments.
 - (b) `deep_q_learning_skeleton.py`, file for the deep Q-learning agent, a part of the code is already provided, part still needs to be implemented.

Deliverables

For this assignment you are required to upload a zip-file containing:

1. Code files with your solutions to the coding exercises.
2. A pdf with answers to the questions.

2 Set up

You will need to get a working python3 installation and install open AI gym.

Set up your python environment

For managing python environments, we highly recommend `virtualenv` and `virtualenvwrapper`.

See: <https://virtualenvwrapper.readthedocs.io/en/latest/>

These will allow you to create project-specific python environments. For instance, after installing `virtualenvwrapper`, setting up a project workspace is as simple as

```
mkvirtualenv -p /usr/bin/python3 AIT
workon AIT
pip install ipython #optional if you want ipython
```

Install open AI gym

Now we can install gym and some dependencies. On Debian this would be as follows:

```
sudo apt install gfortran libopenblas-dev liblapack-dev swig
pip install gym[box2d]
```

(For some reason, the box2d environments are not installed by default and hence need to be specified explicitly: <https://github.com/openai/gym/issues/1603>)

Optional: install pytorch

If you want to use the available neural network code

```
pip install torch
```

3 Environments

In this assignment you will work with two environments. The first will be used for the tabular Q-learning part of the assignment, and will be used to develop an understanding of the difficulty of exploration. The second, Lunar Lander, will be used for the deep Q-learning part.

3.1 Q-Learning

For the first part, we will experiment with ‘drunken walk environments’ that are located in `simple_grid.py`. A simple grid environment, completely based on the code of ‘FrozenLake’, credits to the original authors.

The story here is as follows: You are finding your way home (G) after a great party which was happening at (S). Unfortunately, due to recreational intoxication you find yourself only moving into the intended direction 80% of the time, and perpendicular to that the other 20%. To make matters worse, the local community has been cutting the budgets for pavement (.) maintenance, which means that the way to home is full of potholes (H), which are very likely to make you trip. If you fall, you are obviously magically transported back to the party, without getting some of that hard-earned sleep.

There are different maps available, an example is the "walkInThePark" which has the following shape:

```
"walkInThePark": [  
"S.....",  
"....H..",  
".....",  
".....H.",  
".....",  
"...H...G"  
]
```

Another example is the "theAlley" map:

```
"theAlley": [  
"S...H...H...G"  
]
```

An episode ends when the agent trips or when it reaches the goal. Reaching the goal gives a reward of +10. When moving to a pothole tile there is a 20% chance to trip¹, tripping results in a broken leg penalty.

For the specifics of the implementation of this environment, take a look at the source (`simple_grid.py`).

3.2 Deep Q-Learning

We will use the LunarLander-v2 from Gym. Check here for a short description and here for the source code.

¹This replaces the probability to move perpendicular.

4 Part I: Q-Learning

The Q-Learning algorithm allows us to estimate the optimal Q function using only trajectories from the MDP obtained by following some policy.

Q-learning with ϵ -greedy exploration acts in the following way at a timestep t :

1. In the current state, s , take action a such that a is random with probability ϵ and the greedy action ($a = \max_{a \in A} Q(s, a)$) with probability $1 - \epsilon$;
2. Observe the reward and the next state, r and s' .
3. Update the Q-value as follows:

$$Q^{\text{new}}(s, a) = (1 - \alpha)Q^{\text{old}}(s, a) + \alpha[r + \gamma \max_{a' \in A} Q^{\text{old}}(s', a')]$$

Note that when the episode terminates in s' , the update is as follows:

$$Q^{\text{new}}(s, a) = (1 - \alpha)Q^{\text{old}}(s, a) + \alpha r$$

Coding Exercise 1. Implement Q-learning with ϵ -greedy action selection, complete the class given in `q_learning_skeleton.py`.

Question 1. Which environment, "walkInThePark" or "theAlley", is more difficult to learn in? Why?

Walk in the park

We'll start using the "walkInThePark" map.

Question 2. For the "walkInThePark" map, run some experiments for 1000 episodes with the following settings: $\epsilon = 0.05$, $\gamma = 0.9$, $\alpha = 0.1$. Does the agent learn an optimal policy? Why (not)? Report the (greedy) policy that the agent learned.

The alley

Now we will use the "theAlley" map. First you will compute the optimal Q values, Q^* , this can be done by using value iteration.

Question 3. If you were to apply value iteration, how can you deal with the terminal states?

Question 4. Calculate (or compute) Q^* , the optimal Q-values, for the "theAlley" map with $\gamma = 0.9$, `BROKEN_LEG_PENALTY = -10`.

Question 5. Run some experiments for 1000 episodes with the following settings: $\epsilon = 0.05$, $\gamma = 0.9$, $\alpha = 0.1$, `BROKEN_LEG_PENALTY = -10`. Does the agent learn an optimal policy? Why (not)?

Question 6. Now calculate (or compute) Q^* , the optimal Q-values, for the "theAlley" map with $\gamma = 0.9$, `BROKEN_LEG_PENALTY = -5`.

Question 7. Run some experiments for 1000 episodes with the following settings: $\epsilon = 0.05$, $\gamma = 0.9$, $\alpha = 0.1$,

BROKEN_LEG_PENALTY = -5. Does the agent always learn an optimal policy? Why (not)?

Coding Exercise 2. Try to change the exploration strategy of the agent in a way that allows it to find the optimal solution more often (and quicker).

Question 8. Describe and explain your new exploration strategy. Does it help the agent in learn the optimal policy more often/quicker?

5 Part II: Deep Q-Learning

In regular Q-learning, we had a table to look up the Q-value for each state-action pair. In Deep reinforcement learning we instead use function approximation. We define the Q-value as $Q(s, a; \theta)$, where θ are the parameters of the function approximation, in this case a neural network.

In `deep_q_learning_skeleton.py` a basic version of deep Q-learning has already been implemented.

Todo 1. Familiarize yourself with the code in `deep_q_learning_skeleton.py`.

Question 9. Run `deep_q_learning_main.py` a couple of times. What behavior from the agent do you observe? Does it learn to land safely between the flags?

Coding Exercise 3. Complete the class `ReplayMemory` in `deep_q_learning_skeleton.py`. Change `QLearner` so that it uses the experience replay, that is:

1. `store_experience` should be called in the function `process_experience`
2. In `process_experience` sample a batch of "self.batch_size" from the replay memory and update the network using this experience.

Question 10. Again run `deep_q_learning_main.py` a couple of times. What behavior from the agent do you observe? Does it learn to land safely between the flags? Did the agent improve compared to Question 8?

Coding Exercise 4. Now we will use an additional target network, $Q(s, a; \theta^-)$. Initially our Q-network and the target network will have the same parameters. We will use the target network to provide the estimated future values. That is, we change our target from

$$r + \gamma \max_{a'} Q(s', a', \theta) - Q(s, a, \theta)$$

to

$$r + \gamma \max_{a'} Q(s', a', \theta^-) - Q(s, a, \theta)$$

and at the end of every episode we set $\theta^- = \theta$.

1. In `deep_q_learning_main.py` add a target network to the initialization of the `QLearner`.
2. At the end of every episode set $\theta^- = \theta$,

```
self.target_network.load_state_dict(self.Q.state_dict())
```

3. Change `single_Q_update` (and `batch_Q_update`) to use the Q-value estimation for the next state from the target network.

Question 11. Again run `deep_q_learning_main.py` a couple of times. What behavior from the agent do you observe? Does it learn to land safely between the flags? Did the agent improve compared to Question 9?

References

- [1] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dhharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 02 2015.